

KIREPRO1PE

Compression of Programs and the Similarity Distance

Jonas Nim Røssum <jglr@itu.dk>

Supervisor: Mircea Lungu <mlun@itu.dk> **TODO:** Abstract

May 15, 2025

*They took away the old timbers from time to time, and put new and sound ones in their places, so that the vessel became a standing illustration for the philosophers in the mooted question of growth, some declaring that it **remained the same, others that it was not the same vessel.***

— Plutarch, Life of Theseus 23.1

I would like to thank Christian Gram Kalhauge <chrg@dtu.dk> for proposing the idea for this project and for his valuable external supervision and guidance.

Contents

1. Introduction	5
1.1. 1.2 Research Questions	5
1.2. 1.3 Contributions & Paper Organization	5
2. Background & Related Work	7
2.1. Information distance and Similarity Distance Metrics in Software Engineering	7
2.1.1. Lines of Code Changed as a measure of information distance	7
2.1.2. Shortcomings of LoCC	7
2.1.2.1. First problem: Ambiguous LoCC definitions	7
2.1.2.2. Second problem: Rename-detection pitfalls	8
2.1.2.3. Third problem: Automation-driven LoCC spikes	8
2.1.2.4. Fourth problem: Time bias	8
3. Approach (Proposed Metric) - Using Compression Distance to mitigate the problems of LoCC	10
3.1. Mitigating the problem of ambiguous definitions of LoCC and rename-detection pitfalls	10
3.2. Mitigating the problem of automation-driven LoCC spikes	10
3.3. Definition of Compression Distance	10
3.4. Mitigating survivorship bias in compression distance	11
3.5. Normalized Compression Distance (NCD)	11
4. 4 Methodology	13
4.1. 4.1 Data collection	13
4.1.1. 4.1.1 Repository Selection	13
4.1.2. 4.1.2 File Inclusion/Exclusion	13
4.1.3. 4.1.3 API Data retrieval	13
4.2. 4.2 Metric Computation	14
4.2.1. 4.2.1 Concatenated Commit Buffer (CCB) Construction	14
4.2.2. 4.2.2 Compression Setup	14
4.2.3. 4.2.3 CD & Δ CD Calculation	15
4.3. 4.3 Commit Classification	15
4.4. 4.4 Statistical Analysis	15
5. 5 Results	16
5.1. 5.1 RQ1: LoCC vs. Δ CD Correlation	16
5.1.1. 5.1.1 Correlation Results	16
5.2. 5.2 RQ2: Discrimination Across Commit Types	16
5.2.1. 5.2.1 Δ CD Distributions by Category	16
5.3. 5.3 RQ3: Developer Contribution Analysis	17
5.3.1. 5.3.1 Author-Level Aggregates	17
5.3.2. 5.3.3 Observations: Survivorship Bias Skew	19
6. 6 Discussion	20
6.1. 6.1 Interpretation of RQ1 Findings	20
6.2. 6.2 Interpretation of RQ2 Findings	20
6.3. 6.3 Interpretation of RQ3 Findings	20
6.4. 6.5 Practical Implications	20

6.5. 6.5.1 Considerations for End users	20
6.6. 6.6 Limitations	21
6.6.1. Performance:	21
6.6.2. Scalability	21
6.6.3. Biases	21
7. 7 Conclusion & Future Work	22
7.1. 7.1 Summary of Key Contributions	22
7.2. 7.2 Future Work	22
7.2.1. Scalability	22
7.2.2. Performance	22
8. Acknowledgments	23
Bibliography	24
Index of Figures	25
Index of Tables	25
Index of Listings	25

1. Introduction

Software engineering has consistently explored metrics to measure and compare the complexity and evolution of software systems. Being able to quantify the impact of code changes in software systems is central to understanding software evolution, team productivity, and system complexity. There are many ways to quantify the impact of code changes, such as number of issues closed, number of pull requests merged, number of commits, lines of code changed. Among these, *Lines of code changed* (LoCC) [1] is one of the most widely used distance metrics in the software industry [2]. It is often used as a measure of software complexity, maintainability, and productivity.

While LoCC is simple to compute and interpret, it suffers from several key drawbacks (formatting ambiguity, rename-detection pitfalls, automation-driven spikes, and time bias) that can mislead when analyzing developer effort and project activity. In this paper, we will explore the shortcomings of LoCC and propose an alternative metric: Compression Distance (CD), derived from lossless compression algorithms with large search windows. By measuring the compressibility of changes between software revisions, CD offers a novel perspective on software evolution, addressing many of the shortcomings inherent in LoCC.

Through a series of experiments, we evaluate the effectiveness of CD in quantifying code complexity, distinguishing between commit types, and mitigating biases present in LoCC.

The results suggest that CD provides a more nuanced and robust measure of software evolution, paving the way for its adoption in both academic research and industry practice. The results also show that you need to be aware of built in biases in the metric.

1.1. 1.2 Research Questions

This paper investigates three research questions (RQs):

- RQ1: Is the compression distance a more representative metric for quantifying the complexity of a version-controlled software repository than Lines of Code Changed?
- RQ2: To what extent does compression distance discriminate between manual or semi-automatic commit types (e.g., bugfix, feature, refactoring, documentation, style)?
- RQ3: Does compression distance suffer from the same limitations as LoCC in quantifying the contributions of developers?

1.2. 1.3 Contributions & Paper Organization

We make the following contributions:

- We define Compression Distance (CD), a distance metric based on lossless compression, and derive its per-commit delta (ΔCD).
- We implement CD computation as API endpoints in the Git Truck analysis tool, leveraging ZStandard with a 2 MB search window.
- We empirically evaluate CD on two projects (Git Truck and Commitizen), showing varying correlation with LoCC, improved discrimination of commit types, and distinct author-level insights.

The remainder of the paper is organized as follows. Section 2 reviews LoCC and its limitations. Section 3 presents our proposed CD metric and its theoretical foundation. Section 4 details

the methodology: data collection, metric computation, commit classification, and statistical analyses. Section 5 reports results for RQ1-RQ3. Section 6 discusses implications, practical considerations, and limitations. Finally, Section 7 concludes and outlines directions for future work.

2. Background & Related Work

2.1. Information distance and Similarity Distance Metrics in Software Engineering

In the field of information theory, the concept of *information distance* is used to quantify the similarity between two objects [3]. This is done by measuring the amount of information needed to transform one object into another using a mathematical function F . The most common way to measure this distance is by using a *distance metric*, which is a function that quantifies the difference between two objects.

While such a function F only exists in theory, we can still approximate the information distance between two objects using various practical techniques such as diffing and compression algorithms, as we will explore in this paper.

2.1.1. Lines of Code Changed as a measure of information distance

Since many projects utilize version control systems, such as Git [4], for keeping track of changes, we can track the *Lines of code changed* (LoCC) over time using diffing algorithms. The LoCC metric is typically defined as the number of lines added and removed in a commit. This provides a measure of the information distance between revisions of a software system. Git includes this functionality by default using the numstat argument. [5]. This is a commonly used technique used to detect activity in software systems over time [6]. It can be used to assess team velocity, developer productivity and more. These metrics can be automatically obtained via version control systems using tools like Git Truck[7]. LoCC is a useful metric for quantifying contributions or regions of interest in software systems over time and tools like Git Truck have shown the effectiveness of LoCC in the analysis of software evolution [8], [9].

2.1.2. Shortcomings of LoCC

There are however multiple problems when relying on LoCC as a sole metric of productivity.

2.1.2.1. First problem: Ambiguous LoCC definitions

Firstly, the term itself is ambiguous and subjective to the formatting of the code. One could say that line breaks are just an arbitrary formatting character. In reality, the program could have been written in a single line of code.

See the following ambiguous examples in Listing 1. A few questions arise: should these snippets be counted as three separate lines of code or collapsed into a single line? In terms of actual work, a developer who writes either version is equally productive, yet if we simply count physical lines changed, the author of the first listing would be credited with three times the contribution of the second. This discrepancy shows how formatting alone can skew LoCC-based distance measures, as trivial style differences inflate the perceived distance between revisions and undermine the metric's reliability.

<pre>if (foo) { bar(); }</pre>	<pre>if (foo) { bar(); }</pre>	<pre>if (foo) bar();</pre>
----------------------------------	--------------------------------	----------------------------

Listing 1: Ambiguous line counts

2.1.2.2. Second problem: Rename-detection pitfalls

The LoCC metric can be distorted also by file renames, which may artificially inflate contribution counts, since renaming a file requires little effort but appears as significant line changes. Git supports rename detection using a similarity threshold¹, comparing deleted and added files to identify likely renames, utilized by tools such as Git Truck. While Git doesn't exclude renames by default, this tracking can be used to filter them out manually. However, rename detection is inherently ambiguous — at what point should we stop treating a change as a rename and instead count it as having deleted the old file and added a new one?²

The issue is further exacerbated when developers squash commits, potentially losing rename information. Ultimately, it's a trade-off between overcounting trivial renames and missing substantial, legitimate contributions.

2.1.2.3. Third problem: Automation-driven LoCC spikes

Another case where the LoCC metric falls short is when performing automated actions that affects a vast amount of files and leads to spikes in line changes. Examples of this include running formatting and linting scripts or automated lock file updates when installing packages, all of which can lead to astronomical amounts of line changes.

These kinds of changes do not directly reflect genuine development effort, but they still result in high LoCC values from which one often draw this conclusion. If the developers practice good Git hygiene (such as keeping commits small and focused, not using squash merging, and writing descriptive commit or structured commit messages³), you can perform filtering to focus on for example commits fixing bugs, refactoring or implementing features, but in reality, not all teams conform to these practices.

This shows why this metric cannot stand alone as a measure of productivity. This may cause the results to be misinterpreted, overstating the activity levels of certain developers or areas of the system.

2.1.2.4. Fourth problem: Time bias

Tools like Git Truck track the LoCC through the entire history of a project by default, weighing ancient changes as much as recent changes. In the development of Git Truck, this was attempted to be mitigated by looking at blame information⁴ instead and only considering how the files that are still present in the system as of today has changed through time. However, this technique is still prone to the errors introduced by the problem stated in Section 2.1.2.2.

¹<https://git-scm.com/docs/git-log#Documentation/git-log.txt-code-Mltngtcode>^o

²https://en.wikipedia.org/wiki/Ship_of_Theseus^o

³Some projects conform to structured commit messages using tools like Commitizen [10]

⁴<https://github.com/git-truck/git-truck/commit/12582272b5854d6bf23706b292f3519750023fdd>^o

Another attempt to solve this was the introduction of the time range slider⁵, which led the user select a duration of time to analyze and ignore past changes in analysis.

⁵<https://github.com/git-truck/git-truck/pull/731> °

3. Approach (Proposed Metric) - Using Compression Distance to mitigate the problems of LoCC

To mitigate the explored problems with LoCC, we can use a different approach to measure the distance between revisions of a software system. Instead of relying on the number of lines changed, we can use a metric derived from lossless compression algorithms, to measure the distance between revisions.

3.1. Mitigating the problem of ambiguous definitions of LoCC and rename-detection pitfalls

To mitigate the ambiguity described in Section 2.1.2.1, we can instead consider the change in the number of bytes instead of LoCC as a quantifier of developer activity.

However, this means we can no longer rely on line-based diffing algorithms, akin to those used in `TODO:GIT_DIFF_REF` and must use an alternative method to measure the distance between revisions.

We can mitigate the problem with ambiguity and renames by concatenating all the files existing in each commit. We will refer to this as the concatenated commit buffer, CCB. We can measure its size before and after the commit to calculate a **byte distance metric**, see Equation 0.

$$\Delta|R| = |x| - |x - 1|$$

where $\Delta|R|$ is the distance in bytes, $|x|$ is the size of the given commit buffer and $x - 1$ is the size of the previous commit buffer.

This method does not address the problem with automatic actions overstating developer impact.

This leads us to the Compression Distance.

3.2. Mitigating the problem of automation-driven LoCC spikes

Instead of measuring the static distance in bytes before and after the commit, we can instead try to measure the compression of the changes we added. Using a lossless compression algorithm, we compress the concatenation of the given commit buffer x with the newest revision of the project y and compare this value with the one for the previous commit $x - 1$. The hypothesis being that this would make it such that large repetitive actions would have a lower impact, since they would compress better than smaller, but more complex changes. This assumes that the changes added in a commit will compress better in the presence of similar code. This requires the compression algorithm to have search window larger than double the size of the project files we intend to analyze, as explored in [11].

3.3. Definition of Compression Distance

We define the Compression Distance CD metric as a measure of how much a given concatenated commit buffer compresses with the baseline commit buffer:

$$CD(x, y) = |Z(x)| - |Z(x \cup y)| \quad (1)$$

where CD is the compression distance, x is the given concatenated commit buffer and Z is a lossless compression algorithm.

Now we can define the impact of the commit ΔCD , compared to the previous commit buffer, giving us

$$\Delta CD = CD(x) - CD(x - 1) \quad (2)$$

Unlike the normalized compression distance, this metric is not bounded and is dependent on the absolute sizes of the compressed files. This is fine in this scenario, as we want to compare commits in the same project and the magnitude of the compression distance tells us an interesting story about the impact of the commit.

Then, we can compute the change in the size in bytes before and after the commit.

If we attempt to compress the commit buffers in presence of the newest commit buffer of the repository, we get the side effect of introducing survivorship bias into the system. By measuring the compressed distance to the newest commit buffer of the project, we value changes that are more akin to the newest version higher, in order to tell a story about how we got to the newest version and which commits were the most influential in getting there. This might not be what you want, but for some purposes this makes a lot of sense, as long as you keep the survivorship bias in mind. For example, a detour in the project that didn't make it in the newest version is weighted by a negative distance

TODO: Show a commit from twotter (yoink) that has a negative distance

, telling us this brought the project further away from the newest version, but it might still have been a valuable journey to take. Making “mistakes” is what move projects forward, since you might realize what *not* to do, in order to find out what *to do*.

3.4. Mitigating survivorship bias in compression distance

If we intend to value these types of detours better, we can use the magnitude of the distance instead, which then gives us the impact of the commit, no matter if it moved the project closer or further from the project

To mitigate these problems, we can use an alternative approach where we concatenate the entire state of the repository, makes the analysis resistant to renames, as we don't care about file names, we only care about the bytes contained in the commit. However, in this case, we can no longer rely on diffing the concatenated string. We need to use another approach to derive the information distance from before and after a commit.

3.5. Normalized Compression Distance (NCD)

TODO: If also talking about NCD, write about it here

is a way of measuring the similarity between two compressible objects, using lossless compression algorithms such as GZIP and ZStandard. It is a way of approximating the Normalized Information Distance and has widespread uses such as cluster analysis [<https://arxiv.org/abs/cs/0312044>], and it has even been used to train sentiment analysis.

According to [11], NCD is a very good distance measurement, when used in the proper way.

However, nobody has used normalized compression ratio as a distance metric in version controlled systems yet. The hypothesis of this work is that NCD-derived metrics could function as a complement to the more well-known LC and NoC metrics mentioned above, and be resilient to renaming.

4. 4 Methodology

4.1. 4.1 Data collection

During this project, the analysis tools were implemented that were exposed as API endpoints in the Git Truck project. This was due to the foundation for performing analysis of commits were available in this project, to speed up the development of the tool.

We used these endpoints to collect and visualize data about different repositories to draw conclusions about the Compression Distance metric.

We also report the average time to analyze each repository over 10 warm and cold runs, to give an idea of the interactivity of such a tool. The tool caches file buffers when read using git.

4.1.1. 4.1.1 Repository Selection

TODO: Explain time intervals that are meaningful: git truck before hand in, multiple intervals multiple projects. List why chosen

Repository@revision	NEWEST COMMIT		Commits
	Hash	Buffer size (MB)	
Git Truck@ncd [12]	bf46e09	0.365	1356
Git Truck@v2.0.4 [12]	d385ace	0.318	1260
Git Truck@v1.13.0 [12]	71ae30d	0.259	1242
Twoter@main [13]	2a6a407	0.133	234
Commitizen@master [10]	e177141	0.771	1977

Table 1: Git repositories analyzed

From Table 1, we see that all the chosen repositories lie below the maximum of 1KB, which ensures that backreferencing works as intended.

4.1.2. 4.1.2 File Inclusion/Exclusion

We chose to include/exclude certain file extensions, to focus the results on code files.

In general, file extensions commonly associated with code were included, while binary files like images, videos, and audio files as well as miscellaneous files were excluded [14].

If any extensions were found that were neither included nor excluded, an automatic warning was reported in the console, in order to consider whether it should be included or excluded.

4.1.3. 4.1.3 API Data retrieval

The tool is able to go through a specified range of the history of a git repository and compute metrics for each commit.

The current version of the tool is as of writing not published yet and has to be run manually by cloning the source code.

Queries be made using query parameters like so: Among other metrics, the endpoint is able to compute the Compression Distance in relation to the baseline commit, the newest commit at the given revision in the repository.

- Repository can be specified using the `repo=<folder>` parameter⁶
- Baseline branch or revision is specified with the `branch=<revision>` parameter.
- Amount of commits to analyze is specified using the `count = N|Infinity` parameter, going backwards from branch⁷

To generate the data for this project, the following queries were used:

`http://localhost:3000/get-commits/?repo=git-truck&branch=ncd&count=Infinity`⁸

`http://localhost:3000/get-commits/?repo=git-truck&branch=v2.0.4&count=Infinity`⁸

`http://localhost:3000/get-commits/?repo=git-truck&branch=v1.13.0&count=Infinity`⁸

`http://localhost:3000/get-commits/?repo=commitizen&branch=master&count=Infinity`⁸

This produces a CSV output that can be processed in a data processing tool. During this project, Google Sheets was used for this purpose. Useful information and progress updates is reported in the console.

4.2. 4.2 Metric Computation

In this section, we describe the detailed steps to compute Compression Distance (CD) for each commit. The process is divided into four sub-sections:

4.2.1. 4.2.1 Concatenated Commit Buffer (CCB) Construction

We begin by constructing a concatenated commit buffer (CCB) for every commit, reading all the file contents into a single buffer. The contents of the files are obtained via the `git cat-file <blob hash>` command and reading the standard output of the command.

4.2.2. 4.2.2 Compression Setup

Next, we compress each CCB and its paired baseline buffer using ZStandard (zstd) [15].

The ZStandard compression algorithm (zstd) [16], has a window log of 21 [15] for its default 3 level compression, making the window size $2^{21} = 2\text{MB}$. This makes the it suitable for this task, as long as you are aware of the limit and remember to adjust it as needed.

We used the default compression level of 3, which has a window size of $2^{21} B = 2 \text{ MB}$.

We checked that the repositories we used were smaller than half of this size⁸, in order for the compression algorithm to consider the entire commit buffer when attempting to compress it.

⁶The `repo` parameter refers to a specific git repository folder located relative to where the tools was downloaded.

⁷Passing `Infinity` as count makes the tool go through all the commits in the repository.

⁸Due to compressing the newest commit buffer with each commit, meaning we need at least 2x the baseline buffer, assuming that no commit buffer is larger than the final buffer, but there is still some wiggle room for all of the projects, as none of them surpass 1 MB

4.2.3. 4.2.3 CD & Δ CD Calculation

We then calculate the Compression Distance as

$$\Delta\text{CD}(x, y) = |Z(x)| - |Z(x \cup y)|,$$

where Z is a compression function, x is the CBB and y is the baseline CBB.

This method adds intentional survivorship bias into the metric, favoring code that is more easily compressible with the final version.

4.3. 4.3 Commit Classification

Using keyword searching, we were able to categorize the many of the commits automatically. For the Commitizen repository, all the commits were automatically categorized due to the nature of using structured commit messages. We used the following keywords for categorizing the commits:

For Commitizen, we used: Test, Fix, Feat, Refactor, Style, Docs

For Git Truck, we used: Bump, Refactor, Fix, Feature

4.4. 4.4 Statistical Analysis

We evaluated whether per-commit ΔCD aligns with traditional complexity measured by Lines of Code Changed (LoCC). For each repository, we created log-log scatter plots with linear trend-lines and computed the coefficient of determination, R^2 , between ΔCD and LoCC across all commits.

5. 5 Results

5.1. 5.1 RQ1: LoCC vs. Δ CD Correlation

We use the newest version of each project for this analysis

5.1.1. 5.1.1 Correlation Results

PROJECT	REVISION	R^2 , LINEAR	R^2 , POWER
Git Truck	ncd	0.825	0.494
Commitizen	master	0.297	0.506

Table 2: Correlation between LoCC and Δ CD

From Table 2, it seems that for Git Truck, there exists a linear correlation between LoCC and Δ CD, however for Commitizen, a power regression produces a better correlation.

5.2. 5.2 RQ2: Discrimination Across Commit Types

5.2.1. 5.2.1 Δ CD Distributions by Category

For this experiment, we will look at the Commitizen repository. If we plot each category of commit as a series on a log-log scatter-plot, we can see some interesting patterns. See Figure 1.




Figure 1: Log-log scatter-plot of commits in the Commitizen repository, with automatically categorized commits

From the plot, we can see that bug fix commits specifically has a tendency to have a lower Δ CD metric.

We can also see that version bumps vary much less in impact compared to the other categories. Feature commits generally have a larger LoCC than other commits, and might contain more novel code that compress less, compared to bug fixes.

5.3. 5.3 RQ3: Developer Contribution Analysis

For this experiment, we will look at the Git Truck repository across two time periods.

We accumulate the LoCC and ΔCD for each developer throughout the two time periods.

For context, Git Truck was initially developed by a group of four developers. Later, the project was continuously contributed to several developers, and even later Thomas contributed to the project during his master thesis.

We will compare the project before and after the master thesis by Thomas.

5.3.1. 5.3.1 Author-Level Aggregates

See Figure 2 for the distribution over cumulative LoCC and ΔCD over the two time periods.

See Figure 3 for how the cumulative distribution has changes over time.

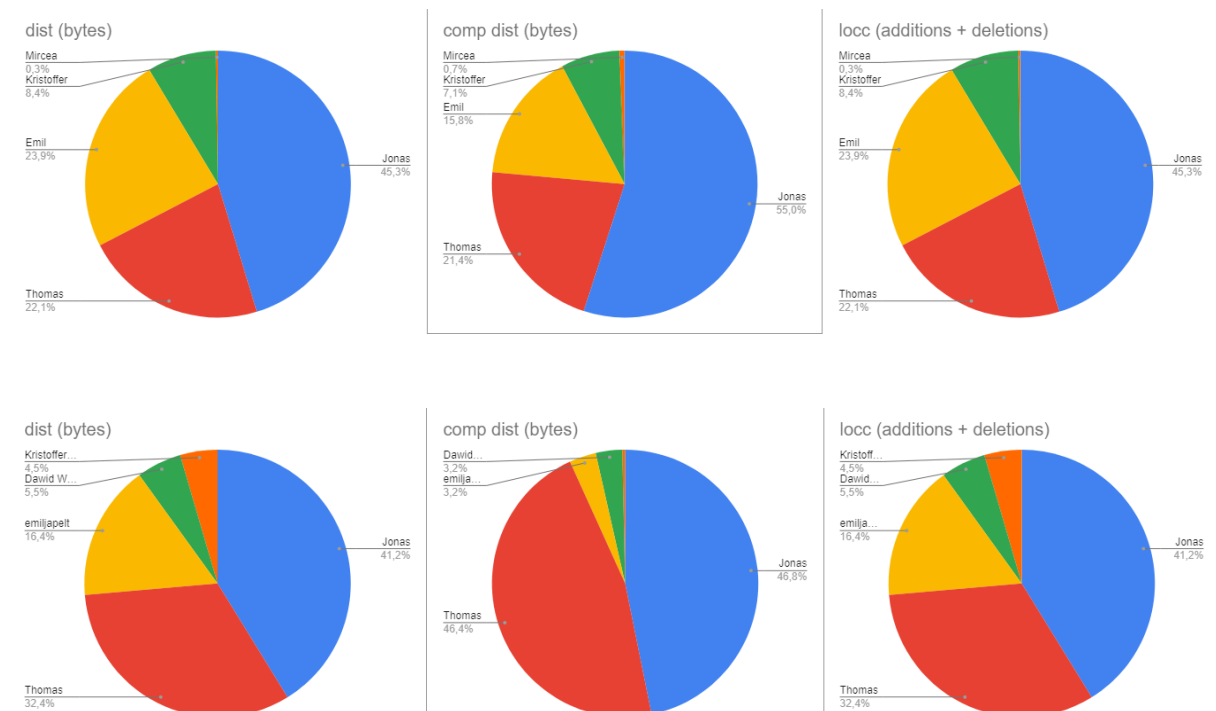


Figure 2: Pie charts of the author distribution before (top) and after (bottom) the thesis project by Thomas









Figure 3: Cumulative area charts of the author distribution over time.
 Overlap (top), stacked (middle) stacked 100% (bottom)
 Δ CD based (left) and LoC based (right)

5.3.2. 5.3.3 Observations: Survivorship Bias Skew

We observe that before Thomas worked intensively on the project, his contribution distributions measured in LoCC and ΔCD were fairly close, however, after his thesis, his share of accumulated ΔCD equalized with that of Jonas.

This is also clear to see on the area time series charts in Figure 3. We can see that using the ΔCD metric, Thomas is able to surpass Jonas in the author distribution. We can also clearly see when Thomas merged his master thesis into the repository and that the merge technique was squashing all of the commits into one, hence the big spike upwards and long period of few commits. This clearly illustrates the information that is lost when developers choose squash merging over regular merges or rebase merging strategies. This is a limitation that neither LoCC, ΔCD or any other metric can preserve, unless the commits are recovered by rebasing them back into the branch.

6. 6 Discussion

6.1. 6.1 Interpretation of RQ1 Findings

In Section 5.1, we observed that there was not necessarily a clear correlation between LoCC and ΔCD , which tells us that the two metrics measure different things and have their own purpose.

6.2. 6.2 Interpretation of RQ2 Findings

In Section 5.2, we observed that certain types of commits contribute more information to the codebase than others. This is intuitive, as introducing a novel feature typically adds more unique content than modifying existing code for bug fixes or version bumps.

6.3. 6.3 Interpretation of RQ3 Findings

There are several hypothesis for the differences observed in . Since Thomas' thesis changed a lot of the codebase, the survivorship bias of the metric favors his recent changes, when judging the history of the project. Another part of the explanation might be that his thesis entailed adding a database to Git Truck, contributing a lot of SQL code to an already TypeScript dominated project, which might also explain the large spike in the ΔCD attribution to Thomas. The hypothesis being that contributing SQL code to the codebase will compress worse than contributing TypeScript code.

This demonstrates the built in survivorship bias that the metric includes and illustrate how you need to be aware of this when using the metric for judging work. For judging work done in group projects, one should select time ranges that correspond to the period of work during the project.

We see that if we measure the author distribution based on cumulative ΔCD contributions, we get survivorship bias built into the metric, which tells us more about who contributed the most to get the system in its current state.

6.4. 6.5 Practical Implications

This study has shown that if implemented with caution for compression window sizes and built in bias, a metric like the compression distance has its place in the arsenal of metrics used in analyzing software evolution. It's a viable supplement to existing metrics and might especially be useful for assessing student projects, as long as detours in the project are also noticed.

6.5. 6.5.1 Considerations for End users

While measuring automatic velocity is a good idea, it is not the only thing to consider. For a user facing project, it is important to also consider the user experience and the impact of changes on the end users. Some small bugfix might have a small impact on the code, but a large impact on the end user.

6.6. 6.6 Limitations

6.6.1. Performance:

We've found that it is much slower to compute than traditional metrics like LoCC built in to Git Truck.

6.6.2. Scalability

For very large repositories, it becomes unfeasible to calculate the compression distance for the entire project at once, as it requires a lot of memory and processing power. This is due to the fact that the compression algorithm needs to keep track of the entire state of the repository in memory. The limited window size [11] leading to distorted distance measurements, especially in worst case when comparing identical objects exceeding these size constraints. This means that we are back to just measuring something similar to the byte distance.

6.6.3. Biases

7. 7 Conclusion & Future Work

7.1. 7.1 Summary of Key Contributions

7.2. 7.2 Future Work

An empirical data study could show how much slower this metric is to compute, compared to existing metrics. It could investigate if compression algorithms in general are slower than diffing algorithms used by LoCC metrics.

7.2.1. Scalability

Attempting to run the tool on a large repository, like the source code for Linux [17] yields a `maxBuffer exceeded` error:

```
RangeError: stdout maxBuffer length exceeded
    at Socket.onChildStdout (node:child_process:481:14)
    at Socket.emit (node:events:507:28)
    at addChunk (node:internal/streams/readable:559:12)
    at readableAddChunkPushByteMode (node:internal/streams/readable:510:3)
    at Socket.Readable.push (node:internal/streams/readable:390:5)
    at Pipe.onStreamRead (node:internal/stream_base_commons:189:23) {
  code: 'ERR_CHILD_PROCESS_STDIO_MAXBUFFER',
  cmd: 'git ls-tree -r HEAD'
```

This shows that future work could be done to investigate whether this metric could scale to very large repositories.

7.2.2. Performance

Future work could be done to investigate methods of speeding up the computational process. The process might be parallelizeable and could utilize using shared memory with dynamic garbage collection, to reduce the memory overhead of the analysis. The current approach caches file buffers, but leaves garbage collection to the runtime.

8. Acknowledgments

The git analysis pipeline from Git Truck was used as a foundation for developing the analysis tool.

Bibliography

- [1] C. to Wikimedia projects, “Source lines of code - Wikipedia.” [Online]. Available: https://en.wikipedia.org/wiki/Source_lines_of_code[◦]
- [2] V. Nguyen, S. Deeds-Rubin, T. Tan, and B. Boehm, “A SLOC counting standard,” in *Cocomo ii forum*, 2007.
- [3] C. to Wikimedia projects, “Information distance - Wikipedia.” [Online]. Available: https://en.wikipedia.org/wiki/Information_distance[◦]
- [4] “Git.” [Online]. Available: <https://git-scm.com/>[◦]
- [5] “Git - git-log Documentation.” [Online]. Available: <https://git-scm.com/docs/git-log#Documentation/git-log.txt-code%E2%80%93numstatcode>[◦]
- [6] M. Goeminne and T. Mens, “Analyzing ecosystems for open source software developer communities,” *Software Ecosystems*. Edward Elgar Publishing, pp. 247–275, 2013.
- [7] K. Højelse, T. Kilbak, J. Røssum, E. Jäpelt, L. Merino, and M. Lungu, “Git-truck: Hierarchy-oriented visualization of git repository evolution,” in *2022 Working Conference on Software Visualization (VISSOFT)*, 2022, pp. 131–140.
- [8] M. Lungu, R.-H. Pfeiffer, M. D’Ambros, M. Lanza, and J. Findahl, “Can git repository visualization support educators in assessing group projects?,” in *2022 Working Conference on Software Visualization (VISSOFT)*, 2022, pp. 187–191.
- [9] A. Neyem, J. Carrasco-Aravena, A. Fernandez-Blanco, and J. P. Sandoval Alcocer, “Exploring the Adaptability and Usefulness of Git-Truck for Assessing Software Capstone Project Development,” in *Proceedings of the 56th ACM Technical Symposium on Computer Science Education V. 1*, 2025, pp. 847–853.
- [10] C. Tools, “commitizen-tools/commitizen: Create committing rules for projects auto bump versions and auto changelog generation.” [Online]. Available: <https://github.com/commitizen-tools/commitizen/>[◦]
- [11] M. Cebrián, M. Alfonseca, and A. Ortega, “Common pitfalls using the normalized compression distance: What to watch out for in a compressor,” 2005.
- [12] G. Truck, “git-truck/git-truck: Git repository visualizations, cumulative contribution statistics and more. Run `npx -y git-truck` to use Git Truck today!.” [Online]. Available: <https://github.com/git-truck/git-truck>[◦]
- [13] T. M. Strings, “themagicstrings/twooter: Twitter replica for DevOps course.” [Online]. Available: <https://github.com/themagicstrings/twooter/>[◦]
- [14] G. Truck, “git-truck/src/routes/get-commits.tsx at ncd · git-truck/git-truck.” [Online]. Available: <https://github.com/git-truck/git-truck/blob/ncd/src/routes/get-commits.tsx#L11-L89>[◦]
- [15] I. Meta Platforms, “zstd/lib/compress/clevels.h at f9938c217da17ec3e9dcd2a2d99c5cf39536aeb9 · facebook/zstd.” [Online]. Available: <https://github.com/facebook/zstd/blob/f9938c217da17ec3e9dcd2a2d99c5cf39536aeb9/lib/compress/clevels.h#L31C7-L31C9>[◦]
- [16] I. Meta Platforms, “facebook/zstd: Zstandard - Fast real-time compression algorithm.” [Online]. Available: <https://github.com/facebook/zstd>[◦]
- [17] torvalds, “torvalds/linux: Linux kernel source tree.” [Online]. Available: <https://github.com/torvalds/linux>[◦]

Index of Figures

Figure 1	Log-log scatter-plot of commits in the Commitizen repository, with automatically categorized commits	16
Figure 2	Pie charts of the author distribution before (top) and after (bottom) the thesis project by Thomas	17
Figure 3	Cumulative area charts of the author distribution over time. Overlap (top), stacked (middle) stacked 100% (bottom) Δ CD based (left) and LoC based (right)	18

Index of Tables

Table 1	Git repositories analyzed	13
Table 2	Correlation between LoCC and Δ CD	16

Index of Listings

Listing 1	Ambiguous line counts	8
-----------	-----------------------------	---